

```
data_path = "gs://us_accidents_projects/US_Accidents_March23.csv"
df = spark.read.csv(data_path, header=True, inferSchema=True)
```

```
df.columns
```

```
['ID',
 'Source',
 'Severity',
 'Start_Time',
 'End_Time',
 'Start_Lat',
 'Start_Lng',
 'End_Lat',
 'End_Lng',
 'Distance(mi)',
 'Description',
 'Street',
 'City',
 'County',
 'State',
 'Zipcode',
 'Country',
 'Timezone',
 'Airport_Code',
 'Weather_Timestamp',
 'Temperature(F)',
 'Wind_Chill(F)',
 'Humidity(%)',
 'Pressure(in)',
 'Visibility(mi)',
 'Wind_Direction',
 'Wind_Speed(mph)',
 'Precipitation(in)',
 'Weather_Condition',
 'Amenity',
 'Bump',
 'Crossing',
 'Give_Way',
 'Junction',
 'No_Exit',
 'Railway',
 'Roundabout',
 'Station',
 'Stop',
 'Traffic_Calming',
 'Traffic_Signal',
```

```
'Turning_Loop',
'Sunrise_Sunset',
'Civil_Twilight',
'Nautical_Twilight',
'Astronomical_Twilight']
```

```
df.describe()
```

```
DataFrame[summary: string, ID: string, Source: string, Severity:
string, Start_Lat: string, Start_Lng: string, End_Lat: string,
End_Lng: string, Distance(mi): string, Description: string, Street:
string, City: string, County: string, State: string, Zipcode: string,
Country: string, Timezone: string, Airport_Code: string,
Temperature(F): string, Wind_Chill(F): string, Humidity(%): string,
Pressure(in): string, Visibility(mi): string, Wind_Direction: string,
Wind_Speed(mph): string, Precipitation(in): string, Weather_Condition:
string, Sunrise_Sunset: string, Civil_Twilight: string,
Nautical_Twilight: string, Astronomical_Twilight: string]
```

```
import pandas as pd_10
print('Columns overview')
pd_10.DataFrame(df.dtypes, columns = ['Column Name', 'Data type'])
```

Columns overview

	Column Name	Data type
0	ID	string
1	Source	string
2	Severity	int
3	Start_Time	timestamp
4	End_Time	timestamp
5	Start_Lat	double
6	Start_Lng	double
7	End_Lat	double
8	End_Lng	double
9	Distance(mi)	double
10	Description	string
11	Street	string
12	City	string
13	County	string
14	State	string
15	Zipcode	string
16	Country	string
17	Timezone	string
18	Airport_Code	string
19	Weather_Timestamp	timestamp
20	Temperature(F)	double
21	Wind_Chill(F)	double
22	Humidity(%)	double

23	Pressure(in)	double
24	Visibility(mi)	double
25	Wind_Direction	string
26	Wind_Speed(mph)	double
27	Precipitation(in)	double
28	Weather_Condition	string
29	Amenity	boolean
30	Bump	boolean
31	Crossing	boolean
32	Give_Way	boolean
33	Junction	boolean
34	No_Exit	boolean
35	Railway	boolean
36	Roundabout	boolean
37	Station	boolean
38	Stop	boolean
39	Traffic_Calming	boolean
40	Traffic_Signal	boolean
41	Turning_Loop	boolean
42	Sunrise_Sunset	string
43	Civil_Twilight	string
44	Nautical_Twilight	string
45	Astronomical_Twilight	string

```
len(df.columns)
```

```
46
```

```
df.printSchema()
```

```
root
```

```
|-- ID: string (nullable = true)
|-- Source: string (nullable = true)
|-- Severity: integer (nullable = true)
|-- Start_Time: timestamp (nullable = true)
|-- End_Time: timestamp (nullable = true)
|-- Start_Lat: double (nullable = true)
|-- Start_Lng: double (nullable = true)
|-- End_Lat: double (nullable = true)
|-- End_Lng: double (nullable = true)
|-- Distance(mi): double (nullable = true)
|-- Description: string (nullable = true)
|-- Street: string (nullable = true)
|-- City: string (nullable = true)
|-- County: string (nullable = true)
|-- State: string (nullable = true)
|-- Zipcode: string (nullable = true)
|-- Country: string (nullable = true)
|-- Timezone: string (nullable = true)
|-- Airport_Code: string (nullable = true)
```

```

|-- Weather_Timestamp: timestamp (nullable = true)
|-- Temperature(F): double (nullable = true)
|-- Wind_Chill(F): double (nullable = true)
|-- Humidity(%): double (nullable = true)
|-- Pressure(in): double (nullable = true)
|-- Visibility(mi): double (nullable = true)
|-- Wind_Direction: string (nullable = true)
|-- Wind_Speed(mph): double (nullable = true)
|-- Precipitation(in): double (nullable = true)
|-- Weather_Condition: string (nullable = true)
|-- Amenity: boolean (nullable = true)
|-- Bump: boolean (nullable = true)
|-- Crossing: boolean (nullable = true)
|-- Give_Way: boolean (nullable = true)
|-- Junction: boolean (nullable = true)
|-- No_Exit: boolean (nullable = true)
|-- Railway: boolean (nullable = true)
|-- Roundabout: boolean (nullable = true)
|-- Station: boolean (nullable = true)
|-- Stop: boolean (nullable = true)
|-- Traffic_Calming: boolean (nullable = true)
|-- Traffic_Signal: boolean (nullable = true)
|-- Turning_Loop: boolean (nullable = true)
|-- Sunrise_Sunset: string (nullable = true)
|-- Civil_Twilight: string (nullable = true)
|-- Nautical_Twilight: string (nullable = true)
|-- Astronomical_Twilight: string (nullable = true)

```

```

print(f'Out of total {df.count()} rows, Displaying first 2 data
rows:')
df.limit(2).toPandas()

```

Out of total 7728394 rows, Displaying first 2 data rows:

23/11/30 02:21:38 WARN package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

	ID	Source	Severity		Start_Time	End_Time
Start_Lat \						
0	A-1	Source2	3	2016-02-08 05:46:00	2016-02-08 11:00:00	
39.865147						
1	A-2	Source2	2	2016-02-08 06:07:59	2016-02-08 06:37:59	
39.928059						

	Start_Lng	End_Lat	End_Lng	Distance(mi)	...	Roundabout	Station
Stop \							
0	-84.058723	NaN	NaN	0.01	...	False	False

```

False
1 -82.831184      NaN      NaN      0.01 ...      False      False
False

    Traffic_Calming Traffic_Signal Turning_Loop Sunrise_Sunset
Civil_Twilight \
0              False              False              False      Night
Night
1              False              False              False      Night
Night

    Nautical_Twilight Astronomical_Twilight
0              Night              Night
1              Night              Day

[2 rows x 46 columns]

```

DATA CLEANING

```

from pyspark import pandas as pd_10

/usr/lib/spark/python/pyspark/pandas/__init__.py:49: UserWarning:
'PYARROW_IGNORE_TIMEZONE' environment variable was not set. It is
required to set this environment variable to '1' in both driver and
executor sides if you use pyarrow>=2.0.0. pandas-on-Spark will set it
for you but it does not work if there is a Spark context already
launched.
    warnings.warn(

from pyspark.sql import functions as ps_10

op = 'Severity'

# find different column types and seggregate to set right defaults
string_cols = [col[0] for col in df.dtypes if col[1] == "string" ]

num_cols = [col[0] for col in df.dtypes if col[1] == "int" or col[1]
== "double" or col[1] == "float" ]

# output could be null
num_cols.remove(op)

bool_cols = [col[0] for col in df.dtypes if col[1] == "boolean"]

print("String columns - ", string_cols)
print("Numeric columns - ", num_cols)
print("Boolean columns - ", bool_cols)

String columns - ['ID', 'Source', 'Description', 'Street', 'City',
'County', 'State', 'Zipcode', 'Country', 'Timezone', 'Airport_Code',
'Wind_Direction', 'Weather_Condition', 'Sunrise_Sunset',

```

```

'Civil_Twilight', 'Nautical_Twilight', 'Astronomical_Twilight']
Numeric columns - ['Start_Lat', 'Start_Lng', 'End_Lat', 'End_Lng',
'Distance(mi)', 'Temperature(F)', 'Wind_Chill(F)', 'Humidity(%)',
'Pressure(in)', 'Visibility(mi)', 'Wind_Speed(mph)',
'Precipitation(in)']
Boolean columns - ['Amenity', 'Bump', 'Crossing', 'Give_Way',
'Junction', 'No_Exit', 'Railway', 'Roundabout', 'Station', 'Stop',
'Traffic_Calming', 'Traffic_Signal', 'Turning_Loop']

```

```

# now we initialize empty rows for each kind of datatype

```

```

df = df.fillna("not_available", string_cols)

```

```

df = df.fillna(0 , num_cols)

```

```

# for bool column, we really can't initialize with 0 or 1

```

```

# Now we count the no of nulls to verify -

```

```

col_vs_nulls = {col:df.filter(ps_10.isnull(df[col[0]))).count() for
col in df.dtypes }
print(col_vs_nulls)

```

```

{('ID', 'string'): 0, ('Source', 'string'): 0, ('Severity', 'int'): 0,
('Start_Time', 'timestamp'): 0, ('End_Time', 'timestamp'): 0,
('Start_Lat', 'double'): 0, ('Start_Lng', 'double'): 0, ('End_Lat',
'double'): 0, ('End_Lng', 'double'): 0, ('Distance(mi)', 'double'): 0,
('Description', 'string'): 0, ('Street', 'string'): 0, ('City',
'string'): 0, ('County', 'string'): 0, ('State', 'string'): 0,
('Zipcode', 'string'): 0, ('Country', 'string'): 0, ('Timezone',
'string'): 0, ('Airport_Code', 'string'): 0, ('Weather_Timestamp',
'timestamp'): 120228, ('Temperature(F)', 'double'): 0,
('Wind_Chill(F)', 'double'): 0, ('Humidity(%)', 'double'): 0,
('Pressure(in)', 'double'): 0, ('Visibility(mi)', 'double'): 0,
('Wind_Direction', 'string'): 0, ('Wind_Speed(mph)', 'double'): 0,
('Precipitation(in)', 'double'): 0, ('Weather_Condition', 'string'):
0, ('Amenity', 'boolean'): 0, ('Bump', 'boolean'): 0, ('Crossing',
'boolean'): 0, ('Give_Way', 'boolean'): 0, ('Junction', 'boolean'): 0,
('No_Exit', 'boolean'): 0, ('Railway', 'boolean'): 0, ('Roundabout',
'boolean'): 0, ('Station', 'boolean'): 0, ('Stop', 'boolean'): 0,
('Traffic_Calming', 'boolean'): 0, ('Traffic_Signal', 'boolean'): 0,
('Turning_Loop', 'boolean'): 0, ('Sunrise_Sunset', 'string'): 0,
('Civil_Twilight', 'string'): 0, ('Nautical_Twilight', 'string'): 0,
('Astronomical_Twilight', 'string'): 0}

```

```

# col_vs_nulls = {col:df.filter(ps_10.isnan(df[col[0]))).count() for
col in df.dtypes if col[1]!='timestamp' }
# print(col_vs_nulls)

```

```
# pd.isna(df) won't work we need to use pysql functions to find null/ na
```

```
state_map = {'AK': 'Alaska',  
             'AL': 'Alabama',  
             'AR': 'Arkansas',  
             'AS': 'American Samoa',  
             'AZ': 'Arizona',  
             'CA': 'California',  
             'CO': 'Colorado',  
             'CT': 'Connecticut',  
             'DC': 'District of Columbia',  
             'DE': 'Delaware',  
             'FL': 'Florida',  
             'GA': 'Georgia',  
             'GU': 'Guam',  
             'HI': 'Hawaii',  
             'IA': 'Iowa',  
             'ID': 'Idaho',  
             'IL': 'Illinois',  
             'IN': 'Indiana',  
             'KS': 'Kansas',  
             'KY': 'Kentucky',  
             'LA': 'Louisiana',  
             'MA': 'Massachusetts',  
             'MD': 'Maryland',  
             'ME': 'Maine',  
             'MI': 'Michigan',  
             'MN': 'Minnesota',  
             'MO': 'Missouri',  
             'MP': 'Northern Mariana Islands',  
             'MS': 'Mississippi',  
             'MT': 'Montana',  
             'NC': 'North Carolina',  
             'ND': 'North Dakota',  
             'NE': 'Nebraska',  
             'NH': 'New Hampshire',  
             'NJ': 'New Jersey',  
             'NM': 'New Mexico',  
             'NV': 'Nevada',  
             'NY': 'New York',  
             'OH': 'Ohio',  
             'OK': 'Oklahoma',  
             'OR': 'Oregon',  
             'PA': 'Pennsylvania',  
             'PR': 'Puerto Rico',  
             'RI': 'Rhode Island',  
             'SC': 'South Carolina',  
             'SD': 'South Dakota',  
             'TN': 'Tennessee',
```

```
'TX': 'Texas',
'UT': 'Utah',
'VA': 'Virginia',
'VI': 'Virgin Islands',
'VT': 'Vermont',
'WA': 'Washington',
'WI': 'Wisconsin',
'WV': 'West Virginia',
'WY': 'Wyoming'}
```

DATA ANALYSIS

Here we find the US states which has most severe accidents since 2020

```
from pyspark.sql import SparkSession

from pyspark.sql.functions import udf
from pyspark.sql.types import StringType

def map_state_code_to_name(code):
    return state_map.get(code, None) # Returns None if the code is
    not found

map_state_udf = udf(map_state_code_to_name, StringType())

accidents_with_state_names = df.withColumn("StateName",
map_state_udf(df["State"]))

#accidents_with_state_names.select("State", "StateName").show()
unique_states_df =
accidents_with_state_names.select("StateName").distinct()
unique_states_df.show()

[Stage 118:=====> (21
+ 2) / 23]
```

```
+-----+
|      StateName|
+-----+
|          Utah|
|      Minnesota|
|          Ohio|
|         Oregon|
|       Arkansas|
|          Texas|
|   North Dakota|
| Pennsylvania|
|   Connecticut|
|        Nebraska|
```



```

|           Vermont|
|           Nevada|
|       Washington|
|           Illinois|
|           Oklahoma|
| District of Columbia|
|           Delaware|
|       New Mexico|
|       West Virginia|
|           Missouri|
+-----+
only showing top 20 rows

```

```
df = accidents_with_state_names
```

```
states_with_severe_accidents = df.select(df.StateName, df.Severity,
df.Start_Time).where(ps_10.year(df.Start_Time)>"2020").filter(df.Severity > 3).groupBy(df.StateName).count().orderBy("count",
ascending=False)
```


```
states_with_severe_accidents_pd =
states_with_severe_accidents.toPandas()
```

```
print(len(states_with_severe_accidents_pd))
```

```
49
```

```
states_with_severe_accidents_pd[:20].plot(kind='barh', x='StateName',
y='count', color="blue", title='Top 20 States with highest number of
accidents with severity > 3 since 2020')
```

```
<AxesSubplot:title={'center':'Top 20 States with highest number of
accidents with severity > 3 since 2020'}, ylabel='StateName'>
```



```


states_with_severe_accidents_pd["percentage"] =
states_with_severe_accidents_pd["count"] /
states_with_severe_accidents_pd["count"].sum()

import matplotlib.pyplot as plt_10

# Create a pie chart
plt_10.figure(figsize=(15, 15))
plt_10.pie(states_with_severe_accidents_pd[:20]['percentage'],
labels=states_with_severe_accidents_pd[:20]['StateName'],
autopct='%1.1f%%', startangle=140)
plt_10.title('Pie Chart of Percentage Count in Every State')
plt_10.show()

```

Pie Chart of Percentage Count in Every State



The next idea is to understand what weather conditions influence the accidents of high severity

```
weather_condt_affecting_accidents = (
    df.select(df.Weather_Condition, df.Severity, df.Start_Time)
#     .where(ps_10.year(df.Start_Time) > "2016")
    .filter(df.Severity > 3)
    .groupBy(df.Weather_Condition)
    .count()
    .orderBy("count", ascending=False)
)
```

```
weather_condt_affecting_accidents_pd =  
weather_condt_affecting_accidents.toPandas()
```

```
print(weather_condt_affecting_accidents_pd[:20])
```

	Weather_Condition	count
0	Fair	59783
1	Clear	26479
2	Mostly Cloudy	23555
3	Cloudy	22828
4	Partly Cloudy	15700
5	Overcast	13867
6	Light Rain	9360
7	not_available	7429
8	Scattered Clouds	5971
9	Light Snow	4759
10	Fog	2774
11	Rain	2029
12	Haze	1073
13	Fair / Windy	1014
14	Heavy Rain	753
15	Light Drizzle	750
16	Snow	718
17	Cloudy / Windy	455
18	T-Storm	409
19	Mostly Cloudy / Windy	407

```
from matplotlib.cm import get_cmap  
import matplotlib.pyplot as plt  
import pandas as pd
```

```
def draw_bar_chart(data, title: str):  
    # Remove or replace None values  
    data = data[data['Weather_Condition'].notna()] # This line  
removes rows where 'Weather_Condition' is None  
    # data['Weather_Condition'] =  
data['Weather_Condition'].fillna('Unknown') # Alternatively, replace  
None with 'Unknown'  
  
    # Proceed with your existing code  
    weather_accident_top20 = data  
  
    # Create a bar chart with colors based on the 'count' column  
    plt.figure(figsize=(12, 6))  
  
    # colormap  
    cmap = get_cmap('viridis')
```

```

# Normalize 'count' values to be used as color intensities
normalize =
plt.Normalize(vmin=weather_accident_top20['count'].min(),
vmax=weather_accident_top20['count'].max())
colors = cmap(normalize(weather_accident_top20['count']))

# Create the bar chart with colormap
bars = plt.bar(weather_accident_top20['Weather_Condition'],
weather_accident_top20['count'], color=colors, edgecolor='black',
linewidth=1.2)

# Adding colorbar for reference
sm = plt.cm.ScalarMappable(cmap=cmap, norm=normalize)
sm.set_array([])
cbar = plt.colorbar(sm, pad=0.1)
cbar.set_label('Count')

# Adding labels and title
plt.xlabel('Weather Condition')
plt.ylabel('Count')
plt.title(title)
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for
better readability

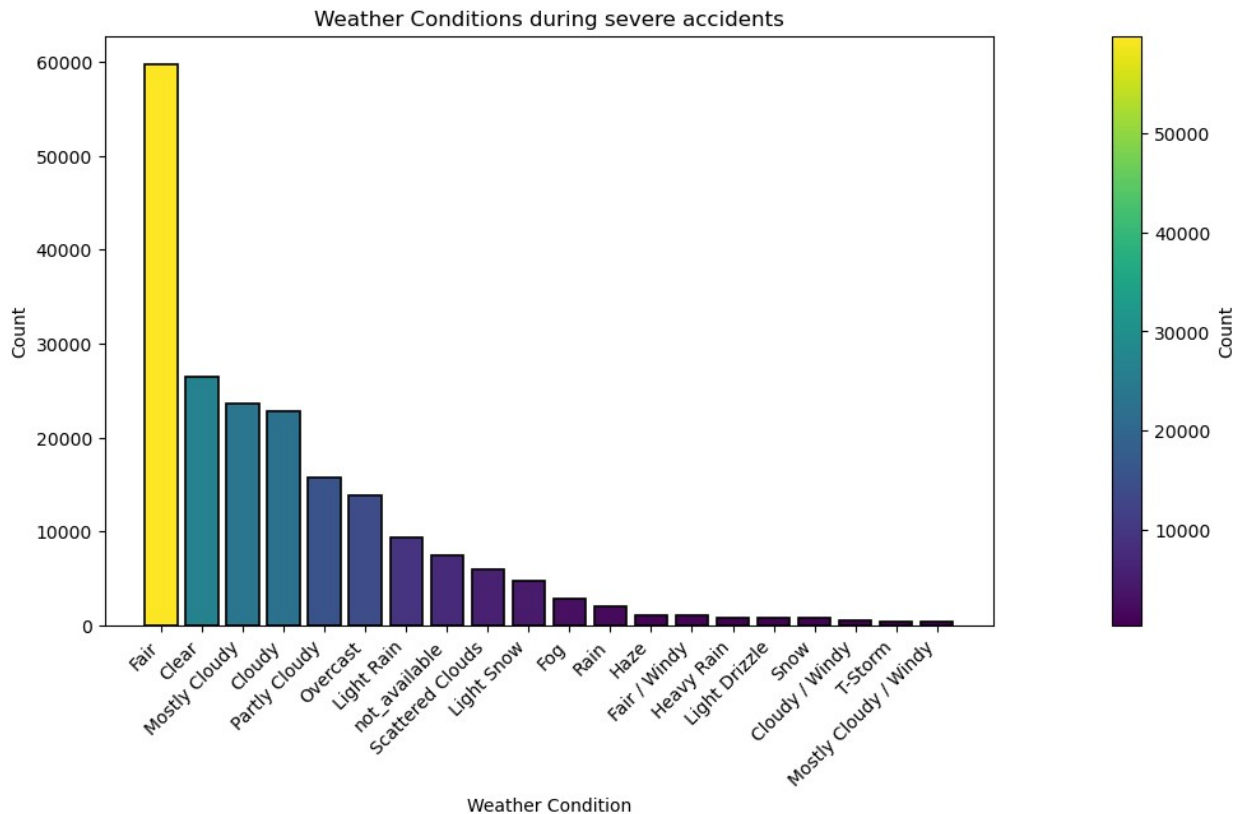
# Display the plot
plt.show()

```

```
print(weather_condt_affecting_accidents_pd[:20].count)
```

		Weather_Condition	count
0	Fair	59783	
1	Clear	26479	
2	Mostly Cloudy	23555	
3	Cloudy	22828	
4	Partly Cloudy	15700	
5	Overcast	13867	
6	Light Rain	9360	
7	not_available	7429	
8	Scattered Clouds	5971	
9	Light Snow	4759	
10	Fog	2774	
11	Rain	2029	
12	Haze	1073	
13	Fair / Windy	1014	
14	Heavy Rain	753	
15	Light Drizzle	750	
16	Snow	718	
17	Cloudy / Windy	455	
18	T-Storm	409	
19	Mostly Cloudy / Windy	407	

```
draw_bar_chart(weather_condt_affecting_accidents_pd[:20],
title="Weather Conditions during severe accidents")
```



```
weather_severity_count = (
    df.select(df.Weather_Condition, df.Severity, df.Start_Time)
      .where(ps_10.year(df.Start_Time) > "2020")
#   .filter(df.Severity > 3)
    .groupBy(df.Weather_Condition, df.Severity)
    .count()
    .orderBy(["Severity","count"], ascending=False)
)

weather_severity_count_pd = weather_severity_count.toPandas()
```

```
print(weather_severity_count_pd[:20])
```

	Weather_Condition	Severity	count
0	Fair	4	37446
1	Cloudy	4	13436
2	Mostly Cloudy	4	8360
3	Partly Cloudy	4	6049
4	not_available	4	3498

5	Light Rain	4	3457
6	Light Snow	4	2008
7	Fog	4	1476
8	Rain	4	756
9	Fair / Windy	4	703
10	Wintry Mix	4	336
11	Haze	4	331
12	Snow	4	324
13	Heavy Rain	4	312
14	Light Drizzle	4	261
15	Cloudy / Windy	4	254
16	T-Storm	4	237
17	Thunder in the Vicinity	4	235
18	Mostly Cloudy / Windy	4	204
19	Light Snow / Windy	4	162

```

top_20_cities = (
    df.select(df.City, df.Start_Time)
      .where(ps_10.year(df.Start_Time) > "2020")
#    .filter(df.Severity > 3)
    .groupBy(df.City)
    .count()
    .orderBy(["count"], ascending=False)
)

```

```
top_20_cities_pd = top_20_cities.toPandas()
```

```
print(top_20_cities_pd)
```

	City	count
0	Miami	128627
1	Orlando	71727
2	Los Angeles	66000
3	Houston	56210
4	Dallas	55086
...
11951	North Greece	1
11952	Leoma	1
11953	Wolsey	1
11954	District 13	1
11955	Red Cloud	1

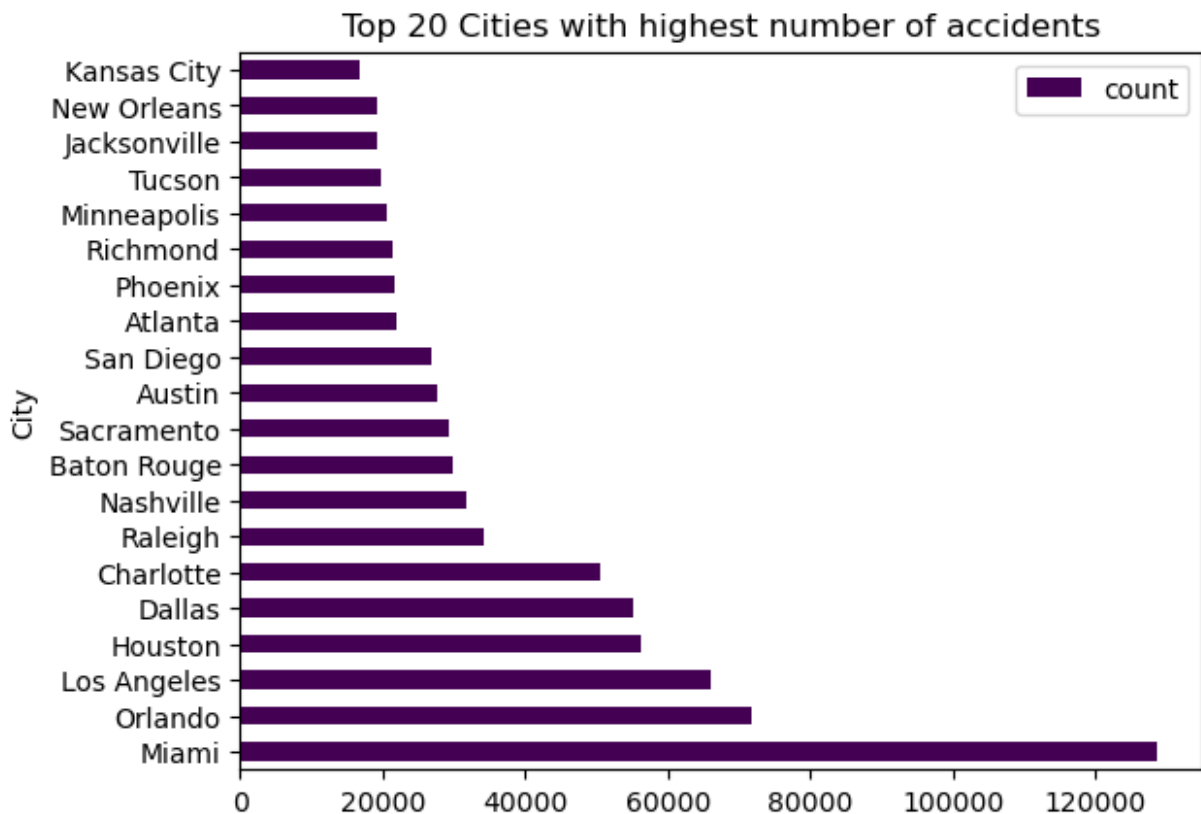
```
[11956 rows x 2 columns]
```

```

top_20_cities_pd[:20].plot(kind='barh', x='City', y='count',
    colormap='viridis', title='Top 20 Cities with highest number of
accidents')

```

```
<AxesSubplot:title={'center':'Top 20 Cities with highest number of accidents'}, ylabel='City'>
```



Miami seems to be the city have the highest number of accidents, (could be because it's a party town:P)

```
df = df.withColumn("Day",
ps_10.dayofmonth(ps_10.to_date(df["Start_Time"])))
df.head()
```

```
Row(ID='A-1', Source='Source2', Severity=3,
Start_Time=datetime.datetime(2016, 2, 8, 5, 46),
End_Time=datetime.datetime(2016, 2, 8, 11, 0), Start_Lat=39.865147,
Start_Lng=-84.058723, End_Lat=0.0, End_Lng=0.0, Distance(mi)=0.01,
Description='Right lane blocked due to accident on I-70 Eastbound at
Exit 41 OH-235 State Route 4.', Street='I-70 E', City='Dayton',
County='Montgomery', State='OH', Zipcode='45424', Country='US',
Timezone='US/Eastern', Airport_Code='KFF0',
Weather_Timestamp=datetime.datetime(2016, 2, 8, 5, 58),
Temperature(F)=36.9, Wind_Chill(F)=0.0, Humidity(%)=91.0,
Pressure(in)=29.68, Visibility(mi)=10.0, Wind_Direction='Calm',
```



```
Wind_Speed(mph)=0.0, Precipitation(in)=0.02, Weather_Condition='Light
Rain', Amenity=False, Bump=False, Crossing=False, Give_Way=False,
Junction=False, No_Exit=False, Railway=False, Roundabout=False,
Station=False, Stop=False, Traffic_Calming=False,
Traffic_Signal=False, Turning_Loop=False, Sunrise_Sunset='Night',
Civil_Twilight='Night', Nautical_Twilight='Night',
Astronomical_Twilight='Night', StateName='Ohio', Day=8)
```

```
# Grouping and aggregating
```

```
daily_accidents = df.groupby("State",
"Day").agg(ps_10.count("*").alias("Accidents"))
```

```
# Calculating the average
```

```
average_accidents =
daily_accidents.groupby("State").agg(ps_10.avg("Accidents").alias("Ave
rage Accidents Per Day"))
```

```
average_accidents.show()
```

```
[Stage 154:=====> (22
+ 1) / 23]
```

```
+-----+-----+
|State|Average Accidents Per Day|
+-----+-----+
| SC | 12340.548387096775 |
| AZ | 5503.5161290322585 |
| LA | 4829.064516129032 |
| MN | 6196.258064516129 |
| NJ | 4539.322580645161 |
| DC | 600.9677419354839 |
| OR | 5795.4838709677415 |
| VA | 9783.90322580645 |
| RI | 547.4516129032259 |
| KY | 1040.4516129032259 |
| WY | 121.19354838709677 |
| NH | 329.4516129032258 |
| MI | 5231.967741935484 |
| NV | 698.8709677419355 |
| WI | 1118.967741935484 |
| ID | 366.96774193548384 |
| CA | 56175.25806451613 |
| CT | 2290.483870967742 |
| NE | 931.2903225806451 |
| MT | 919.2258064516129 |
+-----+-----+
```

```
only showing top 20 rows
```

```
average_accidents_pd = average_accidents.toPandas()

import matplotlib.pyplot as plt_10

# Plot state vs average accidents per day

# Convert the 'Average Accidents Per Day' to a numpy array for easy
manipulation
values = average_accidents_pd['Average Accidents Per Day'].to_numpy()


# Normalize the values to fit the color map
norm = plt_10.Normalize(values.min(), values.max()/3)

# Choose a colormap
cmap = plt_10.cm.inferno

# Apply the colormap
colors = cmap(norm(values))

# Create the plot
plt_10.figure(figsize=(15, 8))
plt_10.bar(average_accidents_pd['State'], values, color=colors)
plt_10.colorbar(plt_10.cm.ScalarMappable(norm=norm, cmap=cmap),
ax=plt_10.gca(), label='Average Accidents Per Day')

plt_10.xlabel('State')
plt_10.ylabel('Average Accidents Per Day')
plt_10.title('Average Road Accidents Per Day in Each State')
plt_10.xticks(rotation=90)
plt_10.show()
```



```
# Sort and take the top 10 states with the highest average accidents
top_states = average_accidents_pd.sort_values(by='Average Accidents
Per Day', ascending=False).head(10)

# geo plot need to do
```

Top 20 dangerous streets for maximum accidents

```
top_20_streets = (
    df.select(df.Street, df.Severity)
      .groupBy(df.Street)
      .count()
      .orderBy(["count"], ascending=False)
)

top_20_streets_pd = top_20_streets.toPandas()


print(len(top_20_streets_pd))
336307
```

```
print(top_20_streets_pd[:20])
```

	Street	count
0	I-95 N	78430
1	I-95 S	74528
2	I-5 N	71968
3	I-10 E	55572
4	I-10 W	53725
5	I-5 S	51781
6	I-80 W	39662
7	I-80 E	36113
8	I-405 N	32364
9	I-75 N	28166
10	I-75 S	27546
11	I-90 E	26426
12	I-15 N	24470
13	I-94 W	24003
14	I-90 W	23279
15	US-101 N	20041
16	I-95	20028
17	I-94 E	19940
18	I-70 E	19697
19	I-15 S	19230

```
top_20_streets_pd[:20].plot(kind='barh', x='Street', y='count',  
colormap='autumn')
```

```
<AxesSubplot:ylabel='Street'>
```



Daily Average per state

```
df = df.withColumn("Day",
ps_10.dayofmonth(ps_10.to_date(df["Start_Time"])))
```

```
df.head()
```

```
Row(ID='A-1', Source='Source2', Severity=3,
Start_Time=datetime.datetime(2016, 2, 8, 5, 46),
End_Time=datetime.datetime(2016, 2, 8, 11, 0), Start_Lat=39.865147,
Start_Lng=-84.058723, End_Lat=0.0, End_Lng=0.0, Distance(mi)=0.01,
Description='Right lane blocked due to accident on I-70 Eastbound at
Exit 41 OH-235 State Route 4.', Street='I-70 E', City='Dayton',
County='Montgomery', State='OH', Zipcode='45424', Country='US',
Timezone='US/Eastern', Airport_Code='KFF0',
Weather_Timestamp=datetime.datetime(2016, 2, 8, 5, 58),
Temperature(F)=36.9, Wind_Chill(F)=0.0, Humidity(%)=91.0,
Pressure(in)=29.68, Visibility(mi)=10.0, Wind_Direction='Calm',
Wind_Speed(mph)=0.0, Precipitation(in)=0.02, Weather_Condition='Light
Rain', Amenity=False, Bump=False, Crossing=False, Give_Way=False,
Junction=False, No_Exit=False, Railway=False, Roundabout=False,
Station=False, Stop=False, Traffic_Calming=False,
Traffic_Signal=False, Turning_Loop=False, Sunrise_Sunset='Night',
```

```
Civil_Twilight='Night', Nautical_Twilight='Night',  
Astronomical_Twilight='Night', StateName='Ohio', Day=8)
```

```
# Grouping and aggregating
```

```
daily_accidents = df.groupBy("State",  
"Day").agg(ps_10.count("*").alias("Accidents"))
```

```
# Calculating the average
```

```
average_accidents =  
daily_accidents.groupBy("State").agg(ps_10.avg("Accidents").alias("Average Accidents Per Day"))
```

```
# Display result
```

```
average_accidents.show()
```

```
[Stage 175:=====> (22  
+ 1) / 23]
```

```
+-----+-----+  
|State|Average Accidents Per Day|  
+-----+-----+  
| SC | 12340.548387096775 |  
| AZ | 5503.5161290322585 |  
| LA | 4829.064516129032 |  
| MN | 6196.258064516129 |  
| NJ | 4539.322580645161 |  
| DC | 600.9677419354839 |  
| OR | 5795.4838709677415 |  
| VA | 9783.90322580645 |  
| RI | 547.4516129032259 |  
| KY | 1040.4516129032259 |  
| WY | 121.19354838709677 |  
| NH | 329.4516129032258 |  
| MI | 5231.967741935484 |  
| NV | 698.8709677419355 |  
| WI | 1118.967741935484 |  
| ID | 366.96774193548384 |  
| CA | 56175.25806451613 |  
| CT | 2290.483870967742 |  
| NE | 931.2903225806451 |  
| MT | 919.2258064516129 |  
+-----+-----+
```

```
only showing top 20 rows
```

```
average_accidents_pd = average_accidents.toPandas()
```

```
import matplotlib.pyplot as plt_10

# Plot state vs average accidents per day

# Convert the 'Average Accidents Per Day' to a numpy array for easy
manipulation
values = average_accidents_pd['Average Accidents Per Day'].to_numpy()

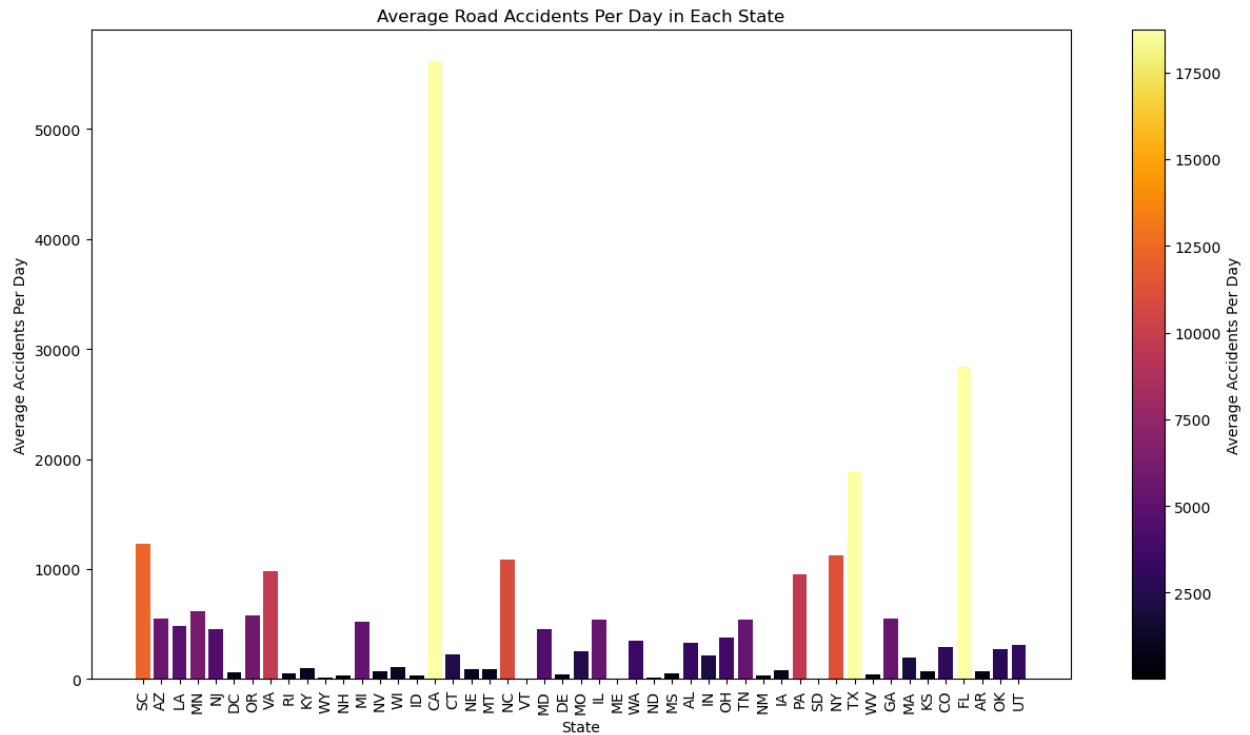
# Normalize the values to fit the color map
norm = plt_10.Normalize(values.min(), values.max()/3)

# colormap - inferno, viridis, plasma
cmap = plt_10.cm.inferno

# Apply the colormap on the normalized values
colors = cmap(norm(values))

# Create the plot
plt_10.figure(figsize=(15, 8))
plt_10.bar(average_accidents_pd['State'], values, color=colors)
plt_10.colorbar(plt_10.cm.ScalarMappable(norm=norm, cmap=cmap),
ax=plt_10.gca(), label='Average Accidents Per Day')

plt_10.xlabel('State')
plt_10.ylabel('Average Accidents Per Day')
plt_10.title('Average Road Accidents Per Day in Each State')
plt_10.xticks(rotation=90)
plt_10.show()
```



```
temp = df.select(df.ID, df.Severity).filter(df.Severity > 3)
temp.collect()[:5]
```

```
[Row(ID='A-620', Severity=4),
 Row(ID='A-1198', Severity=4),
 Row(ID='A-1902', Severity=4),
 Row(ID='A-4144', Severity=4),
 Row(ID='A-4965', Severity=4)]
```